

Sample Input File (students.txt):

```
Alice,85
Bob,67
Charlie,92
David,74
Eva,58
```

Mapper (mapper.py):

```
#!/usr/bin/env python3
import sys

for line in sys.stdin:
    line = line.strip()
    if line:
        name, marks = line.split(',')
        print(f"{name}\t{marks}")
```

Reducer (reducer.py):

```
#!/usr/bin/env python3
import sys

for line in sys.stdin:
    name, marks = line.strip().split('\t')
    marks = int(marks)
    if marks >= 90:
        grade = 'A'
    elif marks >= 80:
        grade = 'B'
    elif marks >= 70:
        grade = 'C'
    elif marks >= 60:
        grade = 'D'
    else:
        grade = 'F'
    print(f"{name}\t{grade}")
```

Execution Commands:

```
# Upload input file to HDFS
hadoop fs -put students.txt /input/students.txt

# Run the MapReduce job
hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar \
-files mapper.py,reducer.py \
-mapper mapper.py \
-reducer reducer.py \
-input /input/students.txt \
-output /output_grades

# View output
hadoop fs -cat /output_grades/part-00000
```

Expected Output:

```
Alice    B
Bob      D
Charlie  A
David    C
Eva      F
```